

Music Genre Classification

A Modular Research Pipeline on GTZAN with Tabular Features and End-to-End Mel Spectrogram Models

Abstract

1. Introduction

2. Problem Analysis

3. Dataset Description

3.1 Repository dataset layout

3.2 Audio metadata audit

3.3 Tabular dataset: feature schema and dimensionality

3.4 Train/validation/test splits (as shipped in this repo)

3.4.1 Tabular split (segment-level, leaky by track)

3.4.2 Mel split (segment-level, *track-disjoint by construction*)

4. Feature Extraction Analysis (with visual evidence)

4.1 Waveform (time-domain amplitude)

4.2 STFT spectrogram (log-frequency, dB)

4.3 Mel spectrogram (dB) — the end-to-end representation

4.4 MFCC (Mel-Frequency Cepstral Coefficients)

4.5 Chromagram (pitch class energy)

4.6 Spectral contrast (peak–valley structure across subbands)

4.7 Tempogram (rhythmic periodicity)

4.8 HPSS (harmonic/percussive source separation)

4.9 Tabular feature space diagnostics (correlation + embeddings)

5. Model Architecture Analysis

5.1 Classical ML baselines (tabular, scikit-learn)

5.2 MLP for tabular features (PyTorch)

5.3 CNN on mel spectrogram (PyTorch)

5.4 LSTM on mel sequences (PyTorch)

6. Training Configuration

6.1 Global settings

6.2 Tabular training (torch MLP)

6.3 Mel feature extraction (end-to-end track)

7. Experimental Results

7.1 Cross-model summary (from evaluation reports)

7.2 End-to-end mel models: CNN vs. LSTM

7.2.1 Training curves

7.2.2 Confusion matrices and ranking curves

7.2.3 Class-wise performance highlights (from evaluation reports)

7.3 Tabular models (with leakage caveat)

7.3.1 LightGBM (best tabular run)

7.3.2 SVM-RBF and XGBoost (competitive tabular runs)

7.3.3 MLP (tabular, neural baseline)

8. Comparative Analysis

8.1 Performance vs. protocol correctness

8.2 Efficiency and complexity trade-offs (as observed from logs)

9. Discussion

9.1 Strengths of the project

9.2 Limitations and risks (evidence-based)

(A) Evaluation leakage in tabular track (confirmed)

(B) Short-context ambiguity for mel models

(C) Dataset size and domain constraints

9.3 Future improvements (actionable)

10. Conclusion

11. References

Music Genre Classification

A Modular Research Pipeline on GTZAN with Tabular Features and End-to-End Mel Spectrogram Models

Authors: *[Your Name]*

Organization: *[University / Lab / Company]*

Date: 2026-05-18

Abstract

This project implements a modular, reproducible pipeline for **music genre classification** on a GTZAN mirror stored under `Data/`. Two modeling tracks are supported: (i) **tabular learning** on pre-extracted 3-second segment features from `Data/features_3_sec.csv` (58-dimensional input after dropping identifier columns), and (ii) **end-to-end neural learning** from raw audio by extracting **mel spectrogram** segments (3 seconds, 128 mel bins) and training deep models (CNN/LSTM). The repository includes stage-wise caching, logging, checkpointing, evaluation reports, and extensive visualizations (audio waveforms, STFT/mel/MFCC/chroma/spectral contrast/tempograms/HPSS, plus embeddings and correlation analysis).

Across the included experimental runs in `outputs/`, the best reported test accuracy is achieved by gradient-boosted tabular models (LightGBM/XGBoost/SVM-RBF) with accuracies around **0.90–0.91**. However, a critical audit of the provided tabular split (`Data/splits/split_v1.csv`) shows severe **track-level leakage: 966/1000 original tracks** have segments appearing in more than one split (train/val/test). Consequently, tabular results should be interpreted as **optimistic upper bounds** under a leaky protocol. In contrast, the mel track uses **group splits by track** with zero overlap across splits, yielding more realistic performance for end-to-end models (CNN: **0.6310**; LSTM: **0.5195**). The report analyzes dataset structure, feature semantics with visual evidence, model architectures and training configurations, and results with a focus on methodological correctness and interpretability.

1. Introduction

Music Genre Classification (MGC) aims to assign a genre label (e.g., blues, classical, hiphop) to an audio recording. It is a canonical task in **Music Information Retrieval (MIR)**, often used to study audio representation learning, signal processing features (STFT, Mel, MFCC), and supervised classification.

Why it matters. Genre recognition supports recommendation systems, music cataloging, playlist generation, and metadata enrichment. Beyond genre, similar pipelines extend naturally to mood, instrument, artist classification, and tagging.

Challenges. - **Inter-class overlap:** many genres share instrumentation or rhythmic patterns (e.g., rock vs. metal). - **Intra-class variability:** genre categories are broad; production style and recording quality vary. - **Temporal structure:** music is non-stationary; short segments may not capture global song identity. - **Evaluation pitfalls:** segmenting tracks and splitting segments naively introduces **data leakage** when segments from the same track appear in train and test.

This repository explicitly focuses on building a **stage-based research pipeline** with caching, reproducibility, visualization, and a “model zoo” spanning classical ML and neural methods.

2. Problem Analysis

MGC is a multi-class classification problem with label space size **C = 10** (confirmed by `Data/features/label_map.json`): `blues`, `classical`, `country`, `disco`, `hiphop`, `jazz`, `metal`, `pop`, `reggae`, `rock` .

Musically, genres often differ along multiple axes: - **Timbre / instrumentation**: MFCC-like spectral envelope patterns often separate classical vs. metal/rock. - **Rhythm / tempo / percussiveness**: hiphop/disco often show strong rhythmic periodicity; classical may be less percussive. - **Harmonic structure**: chroma patterns can be more structured for tonal/harmonic genres; less stable for percussive-heavy segments. - **Spectral brightness**: centroid/rolloff relate to high-frequency content; distorted guitars/hi-hats shift spectral energy upward.

These factors are partially captured by the project's two feature paradigms: - **Tabular summaries** (per-segment statistical aggregates: mean/variance over 3 seconds). - **Time-frequency images/sequences** (mel spectrogram frames over time).

3. Dataset Description

3.1 Repository dataset layout

The repository includes a GTZAN mirror under `Data/` (see `README.md`): -

`Data/genres_original/` : raw audio `.wav` files organized by genre. -

`Data/features_3_sec.csv` : tabular features for ~3-second segments (row-wise samples). -

`Data/features_30_sec.csv` : tabular features for 30-second tracks (track-wise samples). -

`Data/images_original/` : spectrogram images (mirror-dependent). -

`Data/processed/metadata.csv` : metadata generated by the pipeline `preprocess` stage. -

`Data/splits/*` : split definitions used by the pipeline.

3.2 Audio metadata audit

From `Data/processed/metadata.csv` (generated by `src/data/preprocessing.py`): - Total raw audio files scanned: **1000** - Readable/OK: **999** - Unreadable/corrupted: **1**

(`.../jazz/jazz.00054.wav` , marked unreadable) - Sample rate distribution (OK files): **22050**

Hz for all 999 tracks - Duration statistics (OK files): mean $\approx$ **30.02 s** with small variance (min $\approx$ 29.93 s, max $\approx$ 30.65 s)

Figure 1. Raw audio duration distribution (metadata).




Figure 1: Audio duration distribution (sec) from `Data/processed/metadata.csv`.

Figure 2. Raw audio sample rate distribution.




Figure 2: Sample rate distribution (Hz) from `Data/processed/metadata.csv`.

3.3 Tabular dataset: feature schema and dimensionality

`Data/features_3_sec.csv` contains 60 columns, including identifiers and labels: - Identifier columns: `filename`, `length` - Label column: `label` - Feature columns: 57 engineered audio features, including: - `chroma_stft_{mean,var}` - `rms_{mean,var}` - `spectral_centroid_{mean,var}`, `spectral_bandwidth_{mean,var}` - `rolloff_{mean,var}` - `zero_crossing_rate_{mean,var}` - `harmony_{mean,var}`, `perceptr_{mean,var}` (percussive component statistics) - `tempo` - `mfcc1..mfcc20_{mean,var}`

The pipeline's tabular cache builder (`src/data/feature_extraction.py`) drops `filename` and `label`, leaving **58 input dimensions** (57 audio features + `length`).

3.4 Train/validation/test splits (as shipped in this repo)

3.4.1 Tabular split (segment-level, leaky by track)

The provided split file for tabular features is `Data/splits/split_v1.csv` with: - train: **7192** segments - val: **800** segments - test: **1998** segments - total: **9990** segments

Class distribution per split (derived by joining `split_v1.csv` with `features_3_sec.csv` by row index):

label	train	val	test	TOTAL
blues	720	80	200	1000
classical	719	80	199	998
country	718	80	199	997
disco	719	80	200	999
hiphop	718	80	200	998
jazz	720	80	200	1000
metal	720	80	200	1000
pop	720	80	200	1000
reggae	720	80	200	1000
rock	718	80	200	998
TOTAL	7192	800	1998	9990

Critical leakage finding (evidence from filenames). Filenames are of the form `genre.trackid.segment.wav` (e.g., `blues.00000.0.wav`). Auditing `split_v1.csv` shows that **966 out of 1000 track IDs** have segments appearing in more than one split (train/val/test). This means the tabular evaluation protocol is **not track-disjoint**, and test performance is inflated by exposure to the same track's audio characteristics during training.

Consequence: tabular results in `outputs/tabular_*` are best interpreted as performance under a **segment-level split with track leakage**, not as a strict generalization estimate to unseen tracks.

3.4.2 Mel split (segment-level, *track-disjoint by construction*)

The mel pipeline extracts 3-second segments directly from raw audio and assigns a **group ID per original 30-second track** (e.g., `blues.00000`). The split is performed by group (`src/data/split.py::make_group_splits`), and the produced cache `Data/features/mel_features_v1.npz` has: - train: **7183** segments - val: **798** segments - test: **2000** segments - total: **9981** segments

The `groups_train/val/test` sets have **zero overlap** (track-disjoint). Test set is perfectly balanced at 200 segments per class.

	train	val	test	TOTAL
blues	750	50	200	1000
classical	680	118	200	998
country	727	70	200	997
disco	709	90	200	999
hiphop	708	90	200	998
jazz	720	70	200	990
metal	700	100	200	1000
pop	750	50	200	1000
reggae	700	100	200	1000
rock	739	60	200	999
TOTAL	7183	798	2000	9981

4. Feature Extraction Analysis (with visual evidence)

This project supports two complementary notions of “features”: 1. **Tabular engineered features** from `features_3_sec.csv` (statistical summaries over a 3s window). 2. **Time-frequency representations** computed on-demand from raw audio (STFT, mel, MFCC, chroma, etc.) for analysis and for end-to-end learning.


The `visualize` / `visualize-features` stages generate extensive figures under: - `outputs/visualize-features-tabular/figures/` and `outputs/visualize-features-mel/figures/` - `outputs/test-visualize-features/figures/`

Below, each feature is described mathematically and interpreted using actual generated plots.

4.1 Waveform (time-domain amplitude)

Meaning. The waveform is the raw audio signal $x[n]$. While not directly discriminative by itself, it reveals: - amplitude dynamics (compression vs. dynamic range), - silence/structure, - transient density (percussive activity).

Figure 3-6. Example waveforms across genres (3-second windows).



Interpretation. Percussive-heavy segments typically show frequent high-amplitude transients; classical often displays larger macro-dynamics (crescendo/decrescendo) with less dense transient activity. These patterns motivate rhythm- and energy-based features (RMS, tempogram, percussive components).

4.2 STFT spectrogram (log-frequency, dB)


Meaning. The Short-Time Fourier Transform (STFT) reveals energy over time and frequency: $X(m, k) = \sum_{n=0}^{N-1} x[n+mH]w[n]e^{-j2\pi kn/N}$, $S(m, k) = |X(m, k)|$. The plots use

`src/visualization/spectrogram.py::plot_stft` with default `n_fft=2048`, `hop_length=512`.

What it captures. - **Low-frequency structure** (bass lines, kick drum), - **high-frequency content** (hi-hats, distortion, brightness), - transient vs. harmonic texture (vertical vs. horizontal

patterns).

Figure 7-10. STFT examples across genres.




Interpretation. Metal tends to exhibit dense broadband energy (distorted guitars, cymbals), while classical often shows clearer harmonic bands and more separation between instruments. Hiphop can show strong low-frequency energy plus transient percussion patterns. These observations connect to engineered spectral features (centroid, bandwidth, rolloff) and to mel-spectrogram learning.



4.3 Mel spectrogram (dB) — the end-to-end representation

Meaning. The mel spectrogram projects a power spectrogram through a mel filterbank and then applies log compression: $m = 2595 \log(1 + f/700)$, $S\{\text{mel}\} = M S$, $S_{\{\text{dB}\}} = 10 \log(S\{\text{mel}\})$. In this pipeline: - `features.kind=mel_from_audio` extracts mel using `librosa.feature.melspectrogram` and `power_to_db`. - Configuration is in `configs/config.yaml`: - `sample_rate=22050`, `segment_seconds=3.0`, - `n_mels=128`, `n_fft=2048`, `hop_length=512`, - `fmin=20`, `fmax=8000`.

Why useful. Mel spectrograms preserve time–frequency structure and align frequency resolution with human perception, making them a strong input for CNN/LSTM.

Figure 11–14. Mel spectrogram examples.









Interpretation. Relative energy distribution across mel bands (low vs. high) and temporal modulation patterns differentiate genres. For example, sustained harmonic textures appear as continuous horizontal structures; percussive events appear as vertical bursts.

4.4 MFCC (Mel-Frequency Cepstral Coefficients)

Meaning. MFCCs summarize the spectral envelope (timbre) via: 1) mel spectrogram $\log$, 2) Discrete Cosine Transform (DCT). MFCCs are widely used for speech/music timbre characterization and appear as the dominant engineered subfamily in `features_3_sec.csv` (`mfcc1..mfcc20_mean/var`).

Figure 15-18. MFCC heatmaps across genres.









Interpretation. MFCCs compress spectral detail but preserve coarse timbre. Metal/rock timbres often exhibit dense high-frequency presence; classical timbres may show clearer harmonic structure and lower noise-like components. The tabular models heavily rely on MFCC means/variances (40 dimensions out of 58 inputs including **length**), which partly explains strong separability under leaky splitting.

4.5 Chromagram (pitch class energy)

Meaning. A chromagram maps energy to 12 pitch classes (C···B), discarding octave. It reflects harmonic progression and tonality.

Figure 19–22. Chromagrams across genres.







Interpretation. Genres with richer harmonic movement (jazz/classical) often display stronger and more structured chroma activation patterns; highly percussive segments can yield weaker or noisier chroma structure.

4.6 Spectral contrast (peak–valley structure across subbands)

Meaning. Spectral contrast measures energy differences between spectral peaks and valleys across frequency subbands. It is sensitive to instrumentation and timbral “texture”.

Figure 23–26. Spectral contrast examples.







Interpretation. Dense, compressed mixes often reduce contrast in some bands; sparse instrumentations can produce sharper peak–valley differences. This relates to genre-dependent production styles.

4.7 Tempogram (rhythmic periodicity)

Meaning. The tempogram is computed from onset strength over time and captures rhythmic periodicities (tempo-related structure).

Figure 27-30. Tempograms across genres.




Interpretation. Disco/hiphop typically show stronger periodic rhythmic patterns over short windows, whereas classical can be more variable. These properties motivate inclusion of


tempo in the tabular feature set.

4.8 HPSS (harmonic/percussive source separation)

Meaning. HPSS decomposes the STFT into harmonic (sustained) and percussive (transient) components. The tabular feature CSV includes `harmony_*` and `perceptr_*` statistics, and the project provides HPSS visualizations.

Figure 31-34. HPSS visualizations across genres.





Interpretation. Classical/jazz segments often contain strong harmonic structure; hiphop/disco frequently show pronounced percussive energy and transient patterns. The harmonic/percussive balance is a meaningful discriminant for genre texture.

4.9 Tabular feature space diagnostics (correlation + embeddings)

The tabular feature cache is standardized (z-score) on the training split and then visualized via correlation heatmaps and 2D embeddings.

Figure 35. Feature correlation heatmap (subset of first 40 features).






Figure 35: Feature correlation heatmap for tabular inputs.


Figure 36–38. Embeddings of tabular features.

PCA (2D)



t-SNE (2D)





Interpretation. Clear clustering in these embeddings can indicate separability of classes in the engineered feature space. However, because the tabular split is not track-disjoint, apparent separation may partially reflect track-specific signatures leaking across splits.

5. Model Architecture Analysis

The repository includes a “model zoo” spanning classical ML (scikit-learn) and neural models (PyTorch). Model definitions and configuration live in: - `src/models/` (implementations) - `configs/model.yaml` (hyperparameters)

5.1 Classical ML baselines (tabular, scikit-learn)

Implemented via `src/models/sklearn_models.py` with supported names: `knn`, `svm_rbf`, `linear_svm`, `logreg`, `naive_bayes`, `lda`, `qda`, `random_forest`, `extra_trees`, `ada_boost`, `gbdt`, `xgboost`, `lightgbm`.

Rationale. These methods are strong and efficient for low-dimensional tabular data. Ensemble tree methods (ExtraTrees, RandomForest, GBDT, XGBoost, LightGBM) can capture non-linear feature interactions without extensive preprocessing beyond standardization.

Notable implementation details. - `svm_rbf` uses `SVC(probability=True)` to enable ROC/PR curves (computationally heavier but informative). - `linear_svm` uses `LinearSVC`,

which does **not** produce probability estimates; therefore ROC-AUC and PR curves may be unavailable for that run.

5.2 MLP for tabular features (PyTorch)

Defined in `src/models/model.py::MLP` : - Input: 58D vector - Architecture: `Dropout -> (Linear + ReLU + Dropout) × 5 -> Linear(n_classes)` - Hidden sizes: `[512, 256, 128, 64, 32]` , dropout `0.2` (from `configs/model.yaml`)

Strengths/weaknesses. - Flexible non-linear function approximator. - Requires careful regularization and sufficient data; can overfit under leakage or small sample regimes.

5.3 CNN on mel spectrogram (PyTorch)

Defined in `src/models/cnn.py::CNNMel` : - Input: mel tensor shape `(B, 1, n_mels, time)` - Blocks: `Conv2d -> BatchNorm -> ReLU -> MaxPool2d -> Dropout` (repeated 3 times) - Final: `AdaptiveAvgPool2d((1,1)) -> Linear -> logits` - Channels: `[32, 64, 128]` , dropout `0.3`

Rationale. Convolution captures local time–frequency patterns such as rhythmic transients, harmonic stacks, and spectral envelopes.

5.4 LSTM on mel sequences (PyTorch)

Defined in `src/models/lstm.py::LSTMMel` : - Input: sequence shape `(B, time, n_mels)` (mel image reshaped to a time series) - 2-layer bidirectional LSTM with hidden size 128 and dropout 0.2 - Head: `LayerNorm -> Dropout -> Linear(n_classes)`

Rationale. LSTM models temporal dependencies across frames (e.g., evolving rhythm/harmony).

Trade-off. LSTM can be more sensitive to optimization and may underperform a CNN for spectrogram-like inputs when spatial locality in frequency is important.

6. Training Configuration

Configuration is centralized in `configs/config.yaml` and `configs/train.yaml` .

6.1 Global settings

- Seed: `42`
- Device: auto-detect (`cuda > mps > cpu`) with overrides supported (`--device`).
- DataLoader workers: `2`

6.2 Tabular training (torch MLP)

- Epochs: `120`
- Batch size: `256` (train), `512` (eval)

- Optimizer: Adam (`configs/train.yaml`)
- Learning rate: `1.46e-4`
- Weight decay: `1.0e-4`
- Scheduler: ReduceLROnPlateau with factor `0.5` , patience `5` , min LR `1e-6`
- Early stopping: enabled, patience `15` , monitor `val_loss` , mode `min`

6.3 Mel feature extraction (end-to-end track)

- Sampling: `22050 Hz` , mono
- Segment length: `3.0 s`
- STFT: `n_fft=2048` , `hop_length=512`
- Mel bands: `128` , `fmin=20` , `fmax=8000`

7. Experimental Results

All experimental artifacts referenced here are stored under `outputs/` in per-run directories with: - logs: `outputs/<run>/logs/*.log` - figures: `outputs/<run>/figures/*.png` - reports: `outputs/<run>/reports/*`

7.1 Cross-model summary (from evaluation reports)

The table below aggregates metrics from `outputs/*/reports/evaluation_report.md` and approximate training time from the last session in `outputs/*/logs/train.log` .

Run	Accuracy	Macro-F1	ROC-AUC (OvR, macro)	Train time (s)
tabular_lightgbm	0.9129	0.9127	0.9955	13
tabular_svm_rbf	0.9069	0.9068	0.9946	7
tabular_xgboost	0.9044	0.9042	0.9946	16
tabular_extra_trees	0.8904	0.8900	0.9902	3
tabular_knn	0.8744	0.8740	0.9720	6
tabular_random_forest	0.8694	0.8687	0.9887	3
tabular_gbdt	0.8544	0.8549	0.9886	814

Run	Accuracy	Macro-F1	ROC-AUC (OvR, macro)	Train time (s)
tabular_mlp	0.8078	0.8064	0.9777	2092
tabular_qda	0.7487	0.7486	0.9633	0
tabular_logreg	0.7152	0.7129	0.9565	4
tabular_linear_svm	0.6977	0.6919	—	1
tabular_lda	0.6627	0.6611	0.9403	0
mel_cnn	0.6310	0.6211	0.9211	1602
mel_lstm	0.5195	0.5168	0.8782	1185
tabular_naive_bayes	0.4990	0.4795	0.8807	0
tabular_ada_boost	0.4299	0.4155	0.8409	26

Important note (validity). The tabular runs above are evaluated on a segment split with **track leakage** (Section 3.4.1). The mel runs are evaluated on a **track-disjoint** split and should be regarded as a more faithful estimate of generalization to unseen tracks.

7.2 End-to-end mel models: CNN vs. LSTM

7.2.1 Training curves

Figure 39. CNN training curves (outputs/mel_cnn/figures/training_curves.png).




Figure 39: Training curves for mel CNN.

Figure 40. LSTM training curves (outputs/mel_lstm/figures/training_curves.png).






Figure 40: Training curves for mel LSTM.

7.2.2 Confusion matrices and ranking curves

Figure 41–43. CNN evaluation plots.











Figure 44-46. LSTM evaluation plots.

Confusion Matrix (normalized)







7.2.3 Class-wise performance highlights (from evaluation reports)

For `mel_cnn`, the lowest F1 classes are: - rock (F1 ≈ 0.4377), country (≈ 0.4741), pop (≈ 0.4935).

For `mel_lstm`, overall accuracy is lower (0.5195), consistent with the architectural trade-off: treating mel as a 1D time sequence can discard useful local frequency neighborhoods that CNNs exploit.

Observed error patterns (from `outputs/mel_cnn/reports/misclassified.csv`, first 200 errors). The most frequent confusions include: - country $\rightarrow$ classical / rock, - blues $\rightarrow$ country / disco, indicating that short 3-second segments can be ambiguous and that some genre cues require longer temporal context.


7.3 Tabular models (with leakage caveat)


Despite the leaky split, the tabular experiments still reveal which model families best fit the engineered feature space.


7.3.1 LightGBM (best tabular run)

Figure 47–49. LightGBM evaluation plots.

Confusion Matrix (normalized)








Class-wise extremes (from [outputs/tabular_lightgbm/reports/evaluation_report.md](#)). - Strongest: classical (F1 ≈ 0.9630), metal (≈ 0.9602) - Weakest: country (≈ 0.8722), rock (≈ 0.8855)


Observed error patterns (from [outputs/tabular_lightgbm/reports/misclassified.csv](#), first 200 errors). Common confusions include: - country $\rightarrow$ blues, - pop $\rightarrow$ disco / hiphop, - jazz $\rightarrow$ classical, suggesting overlap between harmonic/timbre signatures for short segments.

7.3.2 SVM-RBF and XGBoost (competitive tabular runs)

Figure 50-52. SVM-RBF evaluation plots.

Confusion Matrix (normalized)











Figure 53–55. XGBoost evaluation plots.

Confusion Matrix (normalized)










7.3.3 MLP (tabular, neural baseline)

Figure 56–58. Tabular MLP plots.



Confusion Matrix (normalized)





8. Comparative Analysis

8.1 Performance vs. protocol correctness

The most important comparative result in this repository is not simply “which model has the highest accuracy,” but **which evaluation protocol is valid** for the intended claim (generalization to unseen tracks).

- **Tabular track (as currently split):** very high reported accuracies (≈ 0.91) but invalid as a track-level estimate due to leakage (966/1000 tracks cross splits).
- **Mel track (group split):** lower reported accuracy (CNN 0.6310) but evaluated under a track-disjoint protocol.

Recommendation for fair comparison. To compare tabular and mel models fairly, the tabular split must be changed to be **group-disjoint by track ID** (derivable from `filename`). This is consistent with the intent documented in `README.md` (leakage discussion) and would likely reduce tabular performance substantially.

8.2 Efficiency and complexity trade-offs (as observed from logs)

- Tree ensembles (ExtraTrees/RandomForest) train in seconds and perform strongly under the leaky protocol.
 - Boosted methods (GBDT) can be significantly slower (hundreds of seconds) but remain competitive.
 - Torch MLP and mel CNN/LSTM require minutes of training (per `train.log` timestamps), reflecting higher computational cost.
-

9. Discussion

9.1 Strengths of the project

- **Modular pipeline:** staged execution (`preprocess` , `feature` , `train` , `evaluate` , `visualize`) with structured outputs.
- **Reproducibility:** seed control, cached features, split files, and `config_snapshot.json` per run.
- **Comprehensive visualization:** waveform, STFT/mel/MFCC/chroma/contrast/tempogram/HPSS, plus embeddings and correlation analysis.
- **Model zoo:** broad baseline coverage for tabular learning, plus CNN/LSTM for mel.

9.2 Limitations and risks (evidence-based)

(A) Evaluation leakage in tabular track (confirmed)

The current tabular split file is **not** track-disjoint, strongly inflating performance. This issue is not hypothetical; it is measurable directly from filenames and split assignments.

(B) Short-context ambiguity for mel models

Mel models operate on 3-second segments. Many genre cues are global (song-level arrangement) rather than local (short snippet). Confusions such as country ↔ rock/classical in `mel_cnn` align with this limitation.

(C) Dataset size and domain constraints

GTZAN is relatively small and can be sensitive to recording artifacts and duplication issues across mirrors. The pipeline correctly includes metadata checks and corruption handling (one unreadable file detected).

9.3 Future improvements (actionable)

1. **Fix tabular split to be track-disjoint** by grouping `filename` into track IDs (e.g., `genre.00000`) and splitting by group, mirroring the mel pipeline.
2. **Add a track-level evaluation aggregation** for mel: aggregate segment predictions into a track prediction (majority vote or mean probability), matching real-world “song-level” classification.

3. **Stronger mel architectures:** consider CRNNs (CNN front-end + recurrent temporal modeling) or attention-based pooling over time.
 4. **Augmentation ablations:** the mel extraction code supports an augmentation config (`src/data/augmentation.py`); systematic experiments could improve robustness.
 5. **Calibrated comparisons:** ensure identical track-disjoint splits across all modeling tracks and report confidence intervals via repeated splits.
-

10. Conclusion

This repository provides a well-engineered research pipeline for GTZAN music genre classification, integrating feature caching, logging, checkpointing, and rich visualization. Empirically, the best reported tabular models achieve ≈ 0.91 test accuracy, but this performance is not a valid track-level generalization estimate due to severe segment leakage across splits. The end-to-end mel pipeline, evaluated under a correct track-disjoint split, yields more conservative but methodologically sound results (CNN 0.6310; LSTM 0.5195). The project's most impactful next step is to align the tabular evaluation protocol with the mel pipeline via group-disjoint splitting, enabling fair scientific comparisons between engineered features and learned time-frequency representations.

11. References

1. G. Tzanetakis and P. Cook, "Musical genre classification of audio signals," *IEEE Transactions on Speech and Audio Processing*, 2002. (GTZAN)
2. B. McFee et al., "librosa: Audio and music signal analysis in Python," *Proceedings of the 14th Python in Science Conference*, 2015.
3. F. Pedregosa et al., "Scikit-learn: Machine Learning in Python," *Journal of Machine Learning Research*, 2011.
4. A. Paszke et al., "PyTorch: An Imperative Style, High-Performance Deep Learning Library," *NeurIPS*, 2019.
5. T. Chen and C. Guestrin, "XGBoost: A Scalable Tree Boosting System," *KDD*, 2016.
6. G. Ke et al., "LightGBM: A Highly Efficient Gradient Boosting Decision Tree," *NeurIPS*, 2017.